

An Introduction to Probability

Conditional Probability

Martin Summer

18 January, 2026

Conditional Probability

Conditional Probability

- **Conditional probability** is a basic tool of probability theory.
- Particularly relevant in **Finance** for analyzing dependencies and risk.
- Often obscured by complex terminology despite simple ideas.

Motivation: A Striking Scenario

Motivation: A Striking Scenario

- Imagine evaluating the safety of a **bond portfolio**:
 - Bonds are highly rated and diversified.
 - During a global recession:
 - Defaults occur, unraveling the portfolio.
 - Losses mount unexpectedly.
- Key lesson: **Underestimating event connections** leads to catastrophic risks.
- **Conditional probability** enables us to model dependencies effectively.

Why Conditional Probability Matters

Why Conditional Probability Matters

- Mastering conditional probability is crucial for:
 - Pricing financial instruments.
 - Assessing credit risk.
 - Making informed investment decisions.
- Neglecting it can lead to systemic failures:
 - Example: The 2007-2008 financial crisis.

Case Study: The Financial Crisis (2007-2008)

Case Study: The Financial Crisis (2007-2008)

- Revealed the dangers of assuming independence between events.
- Highlighted failures in **structured finance**:
 - Underestimation of dependencies.
 - Misjudgment of risk profiles.
- Reference: For a comprehensive discussion, see [Tooze 2018].

Understanding Structured Finance

Understanding Structured Finance

- Bonds:

- Financial instruments with fixed payments and default risks.
- Ratings by agencies like Moody's and Standard & Poor's:
 - High grade (AAA, AA) to speculative (BB, B) and default danger (CCC, C).

Rating Category	Moody's	Standard & Poor's
High grade	Aaa	AAA
	Aa	AA
Medium grade	A	A
	Baa	BBB
Speculative grade	Ba	BB
	B	B
Default danger	Caa	CCC
	Ca	CC
	C	C/D

Pooling and Tranching: An Innovation

Pooling and Tranching: An Innovation

- Structured finance:
 - Pools risky assets.
 - Divides cash flows into **tranches** with distinct risk profiles.
 - Enables creation of investment-grade securities from speculative-grade assets.
- Example products:
 - **Mortgage-backed securities (MBS).**
 - Variants using similar financial engineering concepts.

Simplified Example

Simplified Example

- Based on Karl Schmedder's course on probability.
- Illustrates **structured finance** and its relation to probability.
- Develops an intuitive understanding of pooling, tranching, and conditional dependencies.

A Simple Event Tree for One Bond

A Simple Event Tree for One Bond

- Consider a single bond paying €1 at maturity in the future.
 - Default probability: 10% ($P(D) = 0.1$).
 - Non-default probability: 90% ($P(N) = 0.9$).
- Payoff structure:
 - No Default (N): €1.
 - Default (D): €0.

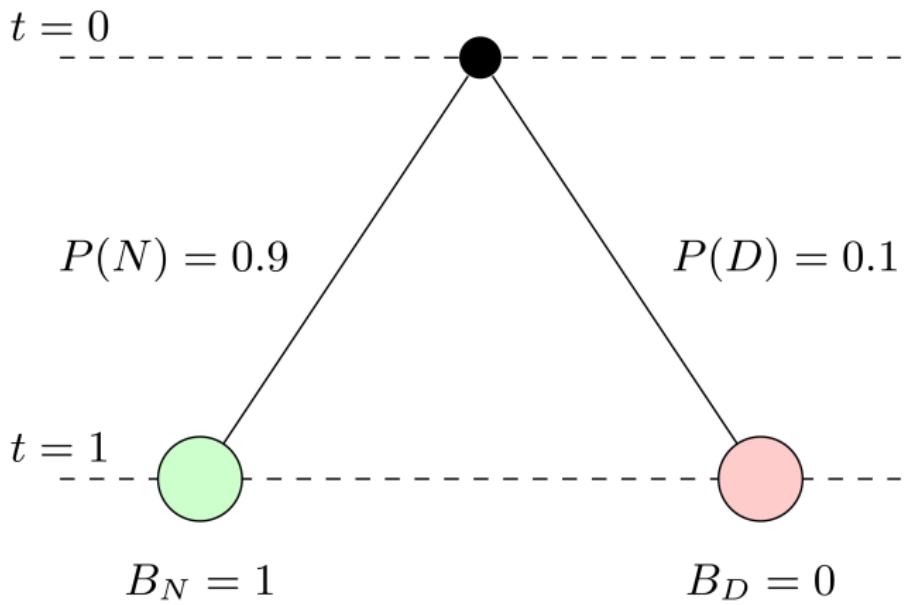


Figure 1

Understanding the Event Tree

Understanding the Event Tree

- Nodes represent states of the bond:
 - $t = 0$: Initial state.
 - $t = 1$: Outcomes (N or D).
- Probabilities:
 - $P(N) = 0.9$.
 - $P(D) = 0.1$.
- Analogy:
 - Coin toss with unequal probabilities:
 - Heads: 90% (No Default).
 - Tails: 10% (Default).
- Probabilistic interpretation:
 - Random experiment with outcomes N and D .

Combining Two Bonds: Independence Assumption

Combining Two Bonds: Independence Assumption

- Portfolio of two bonds.
- **Independence Assumption:**
 - Defaults occur independently.
 - Default of one bond does not influence the other.
- Simplifies calculations:
 - Default probabilities remain uncorrelated.
 - Example: $P(D_1 \cap D_2) = P(D_1) \cdot P(D_2)$.
- Historically justified by:
 - Diversification.
 - Uncorrelated defaults under normal conditions.

Systemic Risks: Challenges to Independence

Systemic Risks: Challenges to Independence

- Systemic risks disrupt independence:
 - Defaults become correlated during crises.
 - Shared macroeconomic factors increase joint defaults.
- Example: 2008 financial crisis:
 - Rising mortgage defaults driven by economic downturn.
 - Increased correlation in bond defaults.
- Critical thinking reveals:
 - Diversification alone cannot guarantee safety.
 - Assumption of independence is fragile.

Event Tree for Two Bonds

Event Tree for Two Bonds

- Two independent bonds:
 - Combine event trees of individual bonds.
 - Visualized as a double event tree.

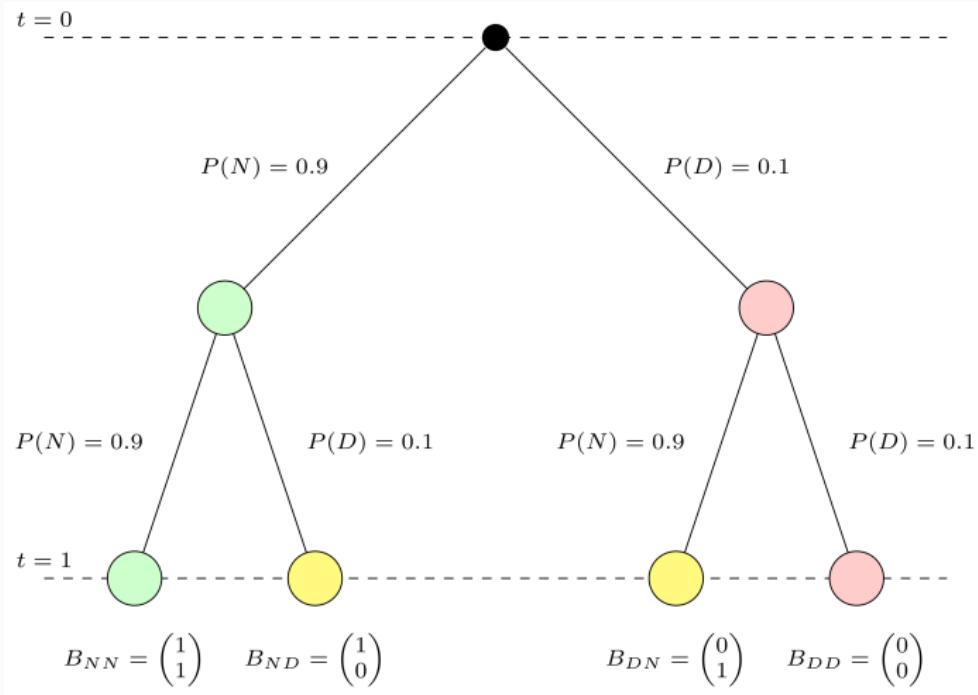


Figure 2

- Outcomes:
 - Each path represents a combination of defaults or no defaults.
 - Example: $B_{NN} = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$ (both bonds pay €1).

Pooling and Tranching: Independent Risks

Pooling and Tranching: Independent Risks

- Re-engineering risk profiles:
 - Pool payoffs of two bonds.
 - Create two new securities:
 1. Pays €1 except when both bonds default.
 2. Pays €0 except when both bonds do not default.
- Probabilities:
 - Both bonds default: $P(D_1 \cap D_2) = 0.1 \cdot 0.1 = 0.01$.
 - Both bonds do not default: $P(N_1 \cap N_2) = 0.9 \cdot 0.9 = 0.81$.

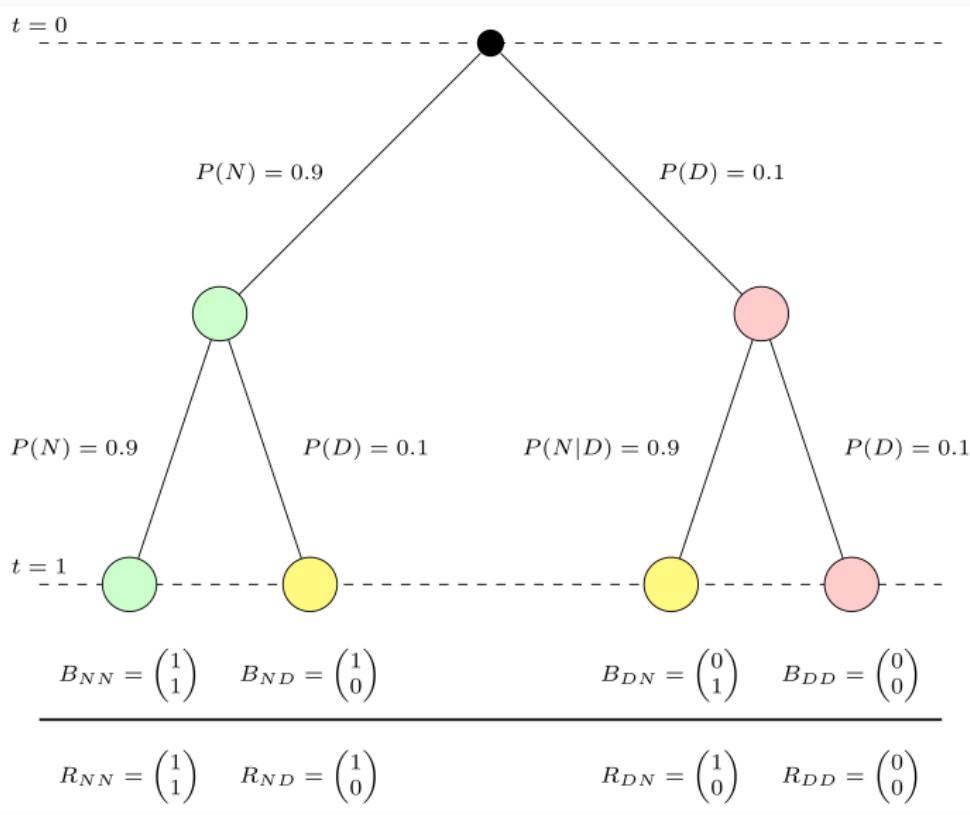


Figure 3

Pooling and Tranching: Dependent Risks

Pooling and Tranching: Dependent Risks

- What happens if independence does not hold?
 - Default probabilities change:
 - $P(D_2 | D_1) = 0.6$.
 - $P(D_2 | N_1) = 0.044$.
 - Increased correlation during systemic events:
 - Shared risks drive joint defaults.

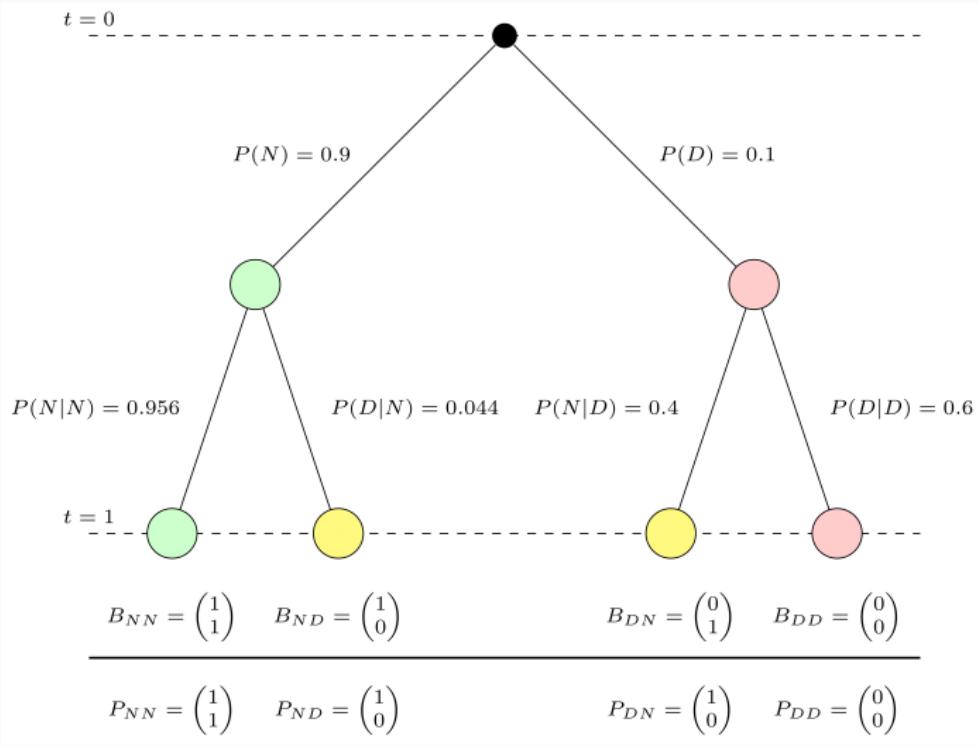


Figure 4

Impact of Dependence on Risk Profiles

Impact of Dependence on Risk Profiles

- Dependent risks change probabilities:
 - Both bonds default: $P(D_1) \cdot P(D_2 | D_1) = 0.1 \cdot 0.6 = 0.06$.
 - Six times higher than under independence.
- Structured finance products fail:
 - Investment-grade tranches lose their safety.
 - Junk plus junk remains junk.

Lessons from Structured Finance

Lessons from Structured Finance

1. Diversification:

- Assets must be from independent sectors.

2. Macroeconomic Stability:

- Low systemic risk is crucial.

3. Transparent Modeling:

- Dependencies must be accounted for.

- Neglecting these led to flawed models and systemic failures during the 2008 crisis.

Conditional Probability

Conditional Probability

- Conditional probability formalizes how the probability of one event changes when another event is known to occur.
- It provides a framework for understanding dependencies quantitatively.

Definition: Conditional Probability

Definition: Conditional Probability



Definition: Conditional Probability

Let B be an event with positive probability. For an arbitrary event A , the **conditional probability** of A given B is:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}, \quad \text{provided } P(B) \neq 0$$

Undefined Conditional Probabilities

Undefined Conditional Probabilities

- Conditional probabilities are **undefined** when the conditioning event B has $P(B) = 0$.
- This distinction is:
 - Irrelevant for **discrete sample spaces**.
 - Crucial in the **general theory**.

Clarifying Conditional Probabilities

Clarifying Conditional Probabilities

- Conditional probabilities represent a **notation** change:
 - Probabilities adjust to reflect known conditions.
- Example: Revisit the financial crisis scenario:
 - Highlighted how dependencies can amplify systemic risk.

The Probability Tree and Conditional Probabilities

The Probability Tree and Conditional Probabilities

- A probability tree is labeled with **edge probabilities**:
 - Representing marginal and conditional probabilities at each level.
 - At $t = 0$
 - $P(B_1 = N) = 0.9, P(B_1 = D) = 0.1$
 - At $t = 1$
 - $P(B_2 = N | B_1 = N) = 0.956,$
 $P(B_2 = D | B_1 = N) = 0.044$
 - $P(B_2 = N | B_1 = D) = 0.4, P(B_2 = D | B_1 = D) = 0.6.$

Defining Probabilities in R

Defining Probabilities in R

- Probabilities defined from the probability tree:

```
# Define the probabilities

# Marginal probabilities for B_1
P_N <- 0.9 # Probability that B_1 does not default
P_D <- 0.1 # Probability that B_1 defaults

# Conditional probabilities for B_2 given B_1
P_N_given_N <- 0.8604/0.9
# B_2 does not default given B_1 does not default
P_D_given_N <- 0.0396/0.9
# B_2 defaults given B_1 does not default
P_N_given_D <- 0.4
# B_2 does not default given B_1 defaults
P_D_given_D <- 0.6
# B_2 defaults given B_1 defaults
```

Computing Joint Probabilities

Computing Joint Probabilities

- Joint probabilities, $P(A \cap B)$, are calculated using the **multiplication rule**:

$$P(A \cap B) = P(A | B) \cdot P(B).$$

```
# Calculate joint probabilities

P_NN <- P_N * P_N_given_N  # Both bonds do not default
P_ND <- P_N * P_D_given_N  # B_1 does not default, B_2 defaults
P_DN <- P_D * P_N_given_D  # B_1 defaults, B_2 does not default
P_DD <- P_D * P_D_given_D  # Both bonds default
```

Simulating Bond Portfolio Defaults

Simulating Bond Portfolio Defaults

Goal:

- Simulate a bond portfolio with two types of bonds (B1 and B2).
- Reproduce a portfolio where default probabilities align with the given contingency table.

Key Concepts:

- Joint probabilities from the contingency table.
- Unconditional and conditional probabilities.

Set up simulation parameters

Set up simulation parameters

```
N <- 5000 # Total number of bonds
P_DD <- 0.06 # P(X = D, Y = D)
P_DN <- 0.04 # P(X = D, Y = N)
P_ND <- 0.04 # P(X = N, Y = D)
P_NN <- 0.86 # P(X = N, Y = N)

# Verify probabilities sum to 1
stopifnot(abs(P_DD + P_DN + P_ND + P_NN - 1) < 1e-6)
```

```
# Simulate joint outcomes based on the contingency table
simulate_defaults <- function(N, probs) {
  sample(
    c("DD", "DN", "ND", "NN"),
    size = N,
    replace = TRUE,
    prob = probs
  )
}
```

```
# Generate portfolio data
portfolio <- data.frame(
  BondID = 1:N,
  BondType = sample(c("B1", "B2"), N,
                    replace = TRUE, prob = c(0.5, 0.5))
)

# Assign joint default outcomes
portfolio$JointOutcome <-
  simulate_defaults(N, c(P_DD, P_DN, P_ND, P_NN))
portfolio$X_Defaulted <-
  portfolio$JointOutcome %in% c("DD", "DN")
portfolio$Y_Defaulted <-
  portfolio$JointOutcome %in% c("DD", "ND")
```

```
# Compute unconditional probabilities
P_X_D <- mean(portfolio$X_Defaulted) # P(X = D)
P_Y_D <- mean(portfolio$Y_Defaulted) # P(Y = D)

# Compute conditional probabilities
P_X_given_Y_D <-
  mean(portfolio$X_Defaulted[portfolio$Y_Defaulted])
# P(X = D | Y = D)
P_Y_given_X_D <-
  mean(portfolio$Y_Defaulted[portfoli0$X_Defaulted])
# P(Y = D | X = D)
```

```
# Display results
cat("Unconditional Probabilities:\n")
cat("P(X = D):", round(P_X_D, 4), "\n")
cat("P(Y = D):", round(P_Y_D, 4), "\n\n")

cat("Conditional Probabilities:\n")
cat("P(X = D | Y = D):", round(P_X_given_Y_D, 4), "\n")
cat("P(Y = D | X = D):", round(P_Y_given_X_D, 4), "\n")
```

```
# Verify calibration matches input probabilities
calibration_check <- table(portfolio$JointOutcome) / N
expected_probs <- c(P_DD, P_DN, P_ND, P_NN)

calibration_result <- data.frame(
  JointOutcome = names(calibration_check),
  Frequency = as.numeric(calibration_check),
  Expected = expected_probs
)
print(calibration_result)
```

Key Takeaways

- **Conditional Probability:**
 - Computed as the relative frequency in a subset of rows where the condition holds.
 - E.g., $P(X = D | Y = D)$ focuses only on rows where `Y_Defaulted` is TRUE.
- **Calibration:**
 - Simulated frequencies align closely with expected probabilities from the contingency table.
- **Practical Application:**
 - Demonstrates how dependency structures in default risks are modeled.

Advanced R Concepts

Advanced R Concepts

In this section, we explore advanced R programming concepts:

1. **Environments**: How R evaluates and stores variables.
2. **Scoping Rules**: How R resolves variable names.
3. **Closures**: Functions that retain the environment where they were created.

Environments

Environments

- An **environment** in R stores objects (variables, functions, etc.).
- The **global environment** stores user-created objects.
- Local variables can override global ones in specific functions.

Example: Global and Local Variables

```
# Global interest rate
interest_rate <- 0.05

# Function to calculate interest
calculate_interest <- function(principal, rate = interest_rate) {
  interest <- principal * rate
  return(interest)
}

# Global calculation
global_interest <- calculate_interest(1000)

# Local override
local_interest <- calculate_interest(1000, rate = 0.07)

cat("Global Interest:", global_interest, "\n")
```

Global Interest: 50

```
cat("Local Interest:", local_interest, "\n")
```

Scoping Rules

Scoping Rules

- R uses **lexical scoping** to find variables:
 - Searches the closest environment first.
 - Moves outward to enclosing environments.

Example: Nested Functions

```
# Global default rates
default_rates <- c(AAA = 0.01, BBB = 0.02, Junk = 0.05)

# Function for conditional default
conditional_default <- function(rating) {
  local_default_rates <- c(
    AAA = uname(default_rates["AAA"]),
    BBB = uname(default_rates["BBB"]),
    Junk = uname(default_rates["Junk"]))
}
return(local_default_rates[rating])
}
```

Lookup Tables

Lookup Tables

- **Lookup tables** map inputs to outputs.
- Example: Default probabilities for credit ratings.

Benefits:

1. Centralizes data for easy updates.
2. Avoids repetitive conditional statements.

Closures

Closures

- **Closures** are functions that retain their creation environment.
- Used to create dynamic and reusable functions.

Example: Probability Calculator Factory

```
# Function factory
probability_calculator_factory <- function(event_probability) {
  function(conditional_probability) {
    joint_probability <- event_probability * conditional_probability
    return(joint_probability)
  }
}

# Create calculators
junk_calculator <- probability_calculator_factory(0.05)
bbb_calculator <- probability_calculator_factory(0.02)

# Calculate joint probabilities
junk_joint <- junk_calculator(0.1)
bbb_joint <- bbb_calculator(0.2)

cat("Joint probability for Junk bonds:", junk_joint, "\n")
```

Joint probability for Junk bonds: 0.005

Analyzing Closures

Analyzing Closures

1. Function Factory:

- Takes `event_probability` as input.
- Returns a function for calculating joint probabilities.

2. Reusable Calculators:

- `junk_calculator` for Junk bonds.
- `bbb_calculator` for BBB bonds.

3. Encapsulation:

- Parameters are “locked in” during function creation.

Bayes' Rule: A Cornerstone of Probability

Bayes' Rule: A Cornerstone of Probability

- Bayes' Rule provides a systematic method for updating probabilities based on new evidence.
- Formula:

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)}$$

- Explanation:
 - A : Observed evidence.
 - B : Hypothesis or prior belief.

Deriving Bayes' Rule

Deriving Bayes' Rule

- From the Multiplication Rule:
 1. $P(B|A)P(A) = P(A \cap B)$
 2. $P(A|B)P(B) = P(A \cap B)$
 - Equating and dividing by $P(A)$:

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)}$$

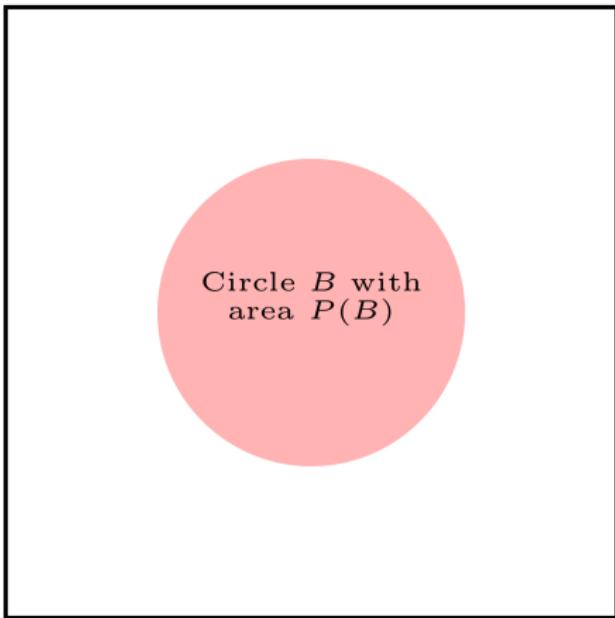
- Significance: Ties prior beliefs to evidence using conditional probabilities.

Speck of Sand: An Intuitive Illustration

Speck of Sand: An Intuitive Illustration

- A square of area 1 represents the sample space.
- Circle B represents the event with area $P(B)$.
- A speck of sand falls randomly in the square.

Square \mathcal{S} with area $P(\mathcal{S}) = 1$

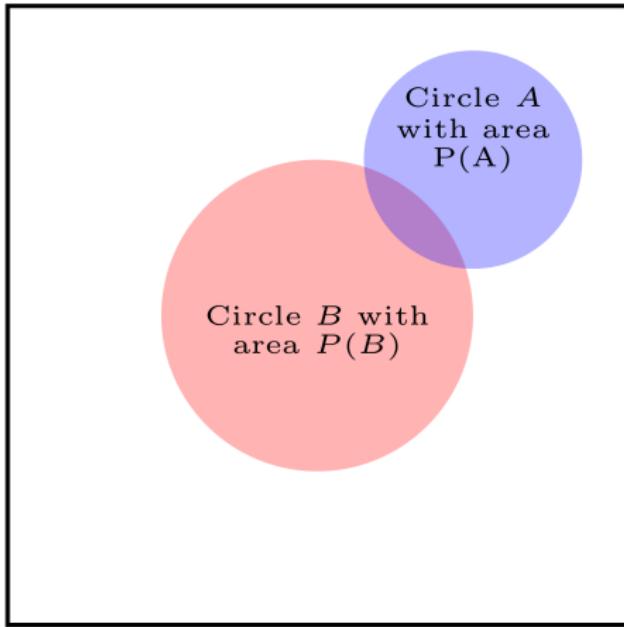


Updating Beliefs

Updating Beliefs

- New Information:
 - The speck is known to land in another circle A inside the square.
- Question:
 - What is $P(B|A)$, the probability that the speck is in B , given it is inside A ?

Square \mathcal{S} with area $P(\mathcal{S}) = 1$



Overlap Between A and B

Overlap Between A and B

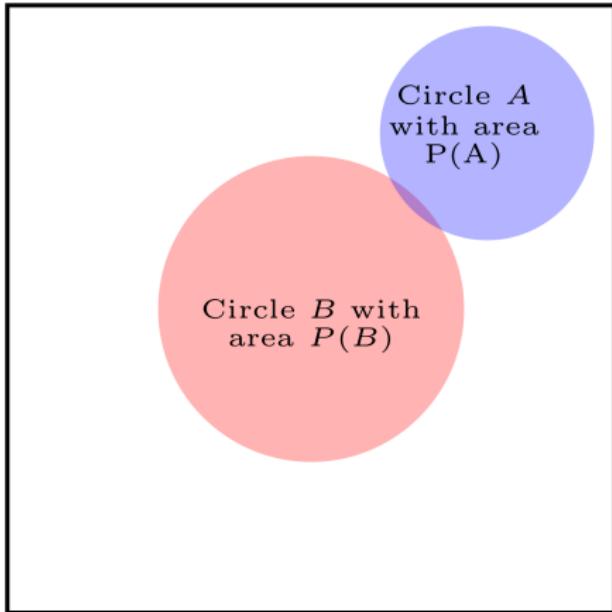
- The updated probability $P(B|A)$ depends on the overlap of B and A :

$$P(B|A) = \frac{\text{Area of } A \cap B}{\text{Area of } A} = \frac{P(A \cap B)}{P(A)}$$

Small Overlap: Low Probability

Small Overlap: Low Probability

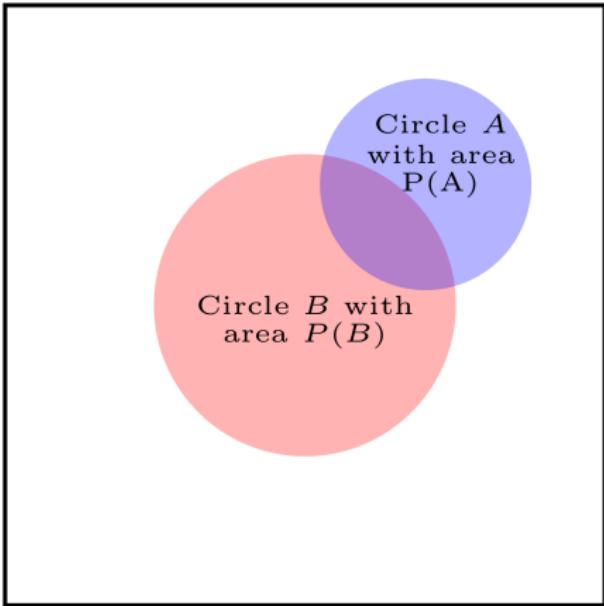
Square \mathcal{S} with area $P(\mathcal{S}) = 1$



Large Overlap: High Probability

Large Overlap: High Probability

Square \mathcal{S} with area $P(\mathcal{S}) = 1$



Bayesian Interpretation

Bayesian Interpretation

- Terms:
 - $P(B)$: Prior probability.
 - $P(A|B)$: Likelihood.
 - $P(A)$: Normalizing constant.
- Process:
 1. Start with prior $P(B)$.
 2. Add evidence $P(A|B)$.
 3. Compute updated belief $P(B|A)$.

Summary

Summary: Probability Concepts

- **Conditional Probability:** $P(A|B) = \frac{P(A \cap B)}{P(B)}$
 - Probability of A given that B has occurred
- **Independence vs. Dependence**
 - Independent: $P(A|B) = P(A)$
 - Dependent: Knowing B changes the probability of A
- **Bayes' Rule:** $P(B|A) = \frac{P(A|B) \cdot P(B)}{P(A)}$
 - Updating beliefs based on new evidence

Summary: Finance Concepts

- Credit Risk and Bond Ratings
 - Default probabilities vary by credit quality
 - Ratings: AAA (safest) to C/D (default)
- Structured Finance: Pooling and Tranching
 - Combine risky assets into pools
 - Create tranches with different risk profiles
- The Independence Trap
 - Assuming independence can drastically underestimate risk
 - Systemic events create correlated defaults
 - Lesson from 2007-2008: dependencies matter!

Summary: R Concepts

- Environments and Scoping
 - Global vs. local variables
 - Lexical scoping: R searches from inner to outer environments
- Lookup Tables
 - Named vectors for mapping inputs to outputs
 - Cleaner than nested if-else statements
- Closures (Function Factories)
 - Functions that “remember” their creation environment
 - Useful for creating specialized calculators

Project: The Coin-Flipping Investment Game

Motivation

- An investor in 2011 believed corporate profit margins would revert to mean
- Year after year, profits stayed elevated while the investor waited
- The mistake: **failing to update beliefs** as disconfirming evidence accumulated
- Winning isn't about having correct priors—it's about **efficiently updating** them

The Game Setup

Two types of coins:

- **Green coins:** 70% Heads, 30% Tails
- **Red coins:** 30% Heads, 70% Tails

Rules:

- Referee secretly draws a coin from a bucket
- You can “buy” each flip at some price
- Heads pays €2.00, Tails pays €0.00

The Key Insight

Worth of a flip depends on the coin:

- Green coin: $0.70 \times 2.00 + 0.30 \times 0.00 = 1.40$ (in €)
- Red coin: $0.30 \times 2.00 + 0.70 \times 0.00 = 0.60$ (in €)

Your task: Estimate which coin is being used to know what to pay

The Challenge

You can't see the coin's color directly. However:

1. Before each round, examine the bucket to estimate the proportion of green vs. red coins (your **prior**)
2. As flips occur, observe the results (your **data**)
3. **Update** your belief about which coin is being used

This is exactly what Bayes' Rule does!

Part 1: Expected Value (Question 1.1)

Suppose you believe there's a 90% chance the coin is green.

- (a) Calculate the expected worth of a single flip.
- (b) Write an R function `calculate_worth()`:

```
# Starter code
calculate_worth <- function(prob_green) {
  # Green coin worth: €1.40
  # Red coin worth: €0.60
  # Your code here
}

# Test: should return €1.32 for prob_green = 0.9
calculate_worth(0.9)
```

Part 2: Bayesian Updating Mathematics

Let G = “coin is green”, D_n = data from n flips with k heads.

$$P(G|D_n) = \frac{P(D_n|G) \cdot P(G)}{P(D_n)}$$

The likelihood follows a binomial distribution:

- $P(k \text{ heads in } n \text{ flips}|G) = \binom{n}{k} (0.7)^k (0.3)^{n-k}$
- $P(k \text{ heads in } n \text{ flips}|R) = \binom{n}{k} (0.3)^k (0.7)^{n-k}$

Part 2: Implementing Bayes' Rule (Question 2.1)

```
update_belief <- function(prior_green, n_heads, n_flips) {  
  p_heads_green <- 0.7  
  p_heads_red <- 0.3  
  
  # Calculate likelihood (hint: use dbinom())  
  likelihood_green <- # Your code  
  likelihood_red <- # Your code  
  
  # Apply Bayes' Rule  
  posterior_green <- # Your code  
  
  return(posterior_green)  
}  
  
# Test: After 10 flips with 2 heads, prior 0.9  
update_belief(0.9, n_heads = 2, n_flips = 10)
```

Part 2: Sequential Updating (Question 2.2)

Write `simulate_round()` that simulates 50 flips and updates beliefs:

```
simulate_round <- function(true_color, prior_green, n_flips = 50) {  
  p_heads <- ifelse(true_color == "green", 0.7, 0.3)  
  
  results <- data.frame(  
    flip_number = 1:n_flips,  
    result = character(n_flips),  
    cumulative_heads = integer(n_flips),  
    posterior_green = numeric(n_flips),  
    worth_estimate = numeric(n_flips)  
  )  
  # Your code: simulate flips and update beliefs  
  return(results)  
}
```

Part 3: Visualizing Overconfidence (Question 3.1)

True coin is **red**, but you start with different priors:

- Prior 0.500: “I have no idea”
- Prior 0.900: “Probably green”
- Prior 0.990: “Almost certainly green”
- Prior 0.999: “Virtually guaranteed green”

```
set.seed(2011) # The year of our hypothetical investor!
priors <- c(0.5, 0.9, 0.99, 0.999)
# Run simulations and plot worth estimates over 50 flips
```

Part 3: The Asymmetry (Question 3.2)

Compare two scenarios:

- When the coin is **red** but you thought it was green
- When the coin is **green** and you thought it was green

Question: Why is overconfidence asymmetrically costly?

(Hint: What do you gain vs. what do you lose?)

Part 4: Profit Calculation (Question 4.1)

You bid your worth estimate on each flip:

- Pay: your worth estimate
- Receive: €2.00 if Heads, €0.00 if Tails

```
calculate_profits <- function(simulation_results) {  
  # For each flip:  
  # - You pay the worth_estimate  
  # - You receive 2.00 if Heads, 0.00 if Tails  
  # - profit = payout - worth_estimate  
  # Return data frame with cumulative_profit added  
}
```

Part 5: Real-World Connection (Question 5.1)

In the coin-flipping game analogy:

- What corresponds to “profit margins are elevated”?
- What corresponds to “profit margins revert to the mean”?
- What corresponds to “assigning a 99.9% prior”?

Key insight: The investor’s mistake wasn’t having wrong initial beliefs, but failing to update efficiently.

Summary: What You'll Learn

1. Bayesian updating provides a principled way to revise beliefs based on evidence
2. Overconfidence is asymmetrically costly: being wrong with high confidence takes longer to correct
3. Evidence required for strong beliefs should be proportional to confidence
4. Efficient updating can compensate for initially incorrect beliefs

Reference

The coin-flipping game analogy is adapted from “Profit Margins, Bayes’ Theorem, and the Dangers of Overconfidence” from the blog *Philosophical Economics* (September 2017).